

# ***k*-Nearest Neighborhood**

**Instructor: Junghye Lee**

**Department of Industrial Engineering**

**[junghyelee@unist.ac.kr](mailto:junghyelee@unist.ac.kr)**

# Simple Analogy


- Tell me about your friends (*who your neighbors are*) and *I will tell you who you are.*




# What is *k*-Nearest Neighborhood (k-NN)?

- A powerful classification algorithm used in pattern recognition.
- *k*-nearest neighbors stores all available cases and classifies new cases based on a ***similarity measure*** (e.g., distance function)
- A **non-parametric** lazy learning algorithm (i.e., instance-based learning, memory-based reasoning, example-based Reasoning)


# $k$ -NN



(a) 1-nearest neighbor



(b) 2-nearest neighbor




(c) 3-nearest neighbor

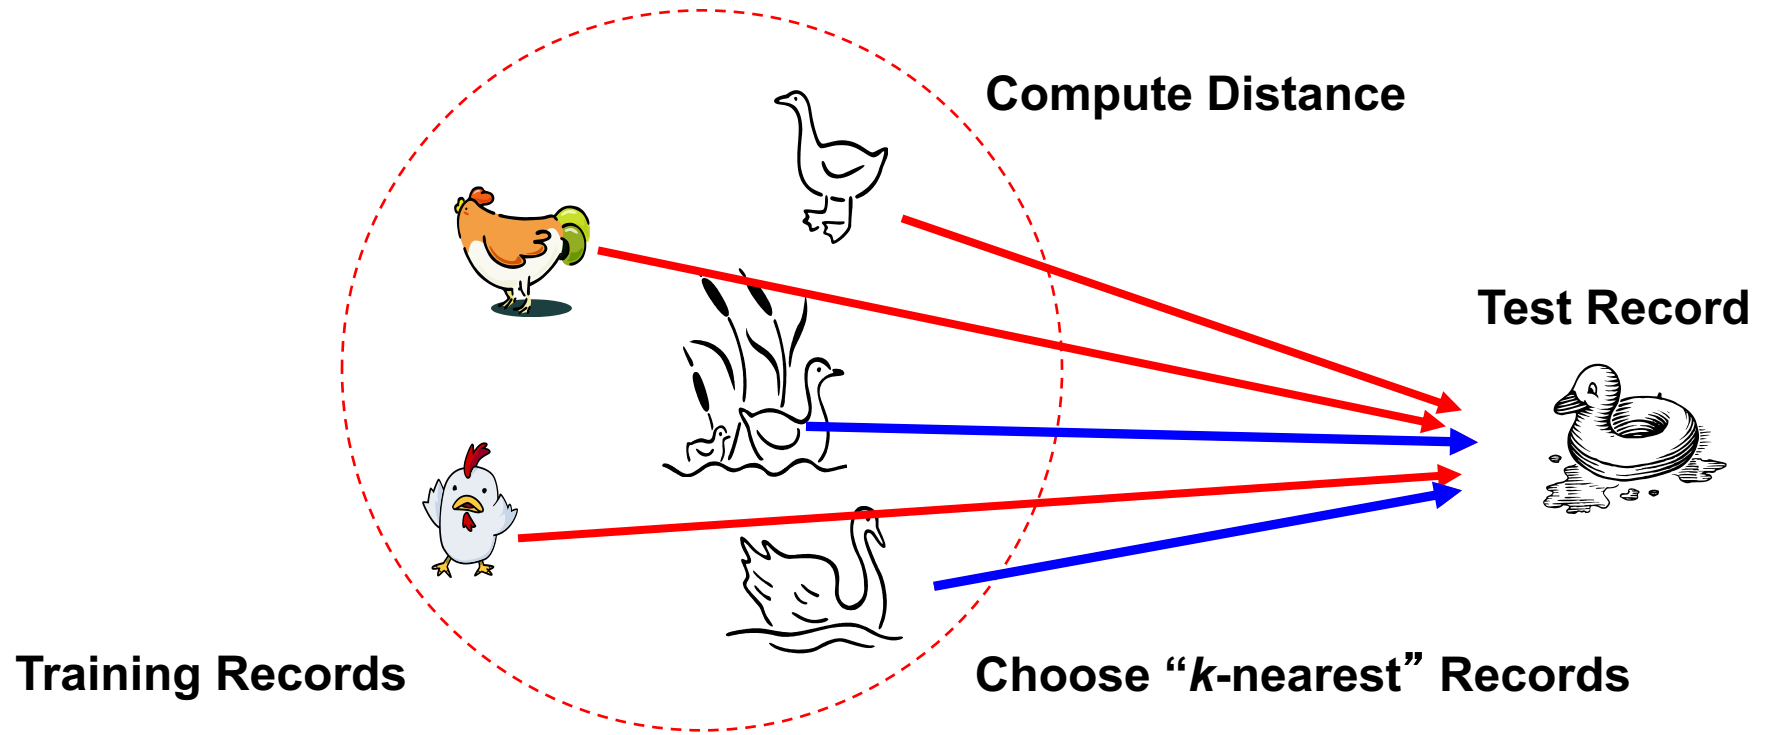
- $k$ -nearest neighbors of a record  $x$  are data points that have the  $k$  smallest distance to  $x$

# *k*-NN: Classification Approach

- An object (a new instance) is classified by a majority votes for its neighbor classes.
- The object is assigned to the most common class amongst its *k*-nearest neighbors (measured by a distance function).



# Distance Measure



# Distance Measure for Continuous Variables

- Euclidean

$$\sqrt{\sum_{i=1}^p (x_i - y_i)^2}$$

- Manhattan

$$\sum_{i=1}^p |x_i - y_i|$$

- Minkowski

$$\left(\sum_{i=1}^p (|x_i - y_i|)^q\right)^{1/q}$$

# Distance between Neighbors

- Calculate the distance between new example (E) and all examples in the training set.
- **Euclidean distance** between two examples.

$$X = [x_1, x_2, x_3, \dots, x_p]$$

$$Y = [y_1, y_2, y_3, \dots, y_p]$$

- The Euclidean distance between  $X$  and  $Y$  is defined as:

$$D(X, Y) = \sqrt{\sum_{i=1}^p (x_i - y_i)^2}$$



# ***k*-Nearest Neighbor Algorithm**


- All the instances correspond to points in an  $p$ -dimensional feature space.
- Each instance is represented with a set of numerical attributes.
- Each of the training data consists of a set of vectors and a class label associated with each vector.
- Classification is done by comparing feature vectors of different  $k$ -nearest points.
- Select the  $k$ -nearest examples to  $E$  in the training set.
- Assign  $E$  to the most common class among its  $k$ -nearest neighbors.

# 3-NN: Example

Customer	Age	Income	No. credit cards	Class	Distance from John
George	35	35K	3	No	$\text{sqrt} [(35-37)^2+(35-50)^2+(3-2)^2]=15.16$
Rachel	22	50K	2	Yes	$\text{sqrt} [(22-37)^2+(50-50)^2+(2-2)^2]=15$
Steve	63	200K	1	No	$\text{sqrt} [(63-37)^2+(200-50)^2+(1-2)^2]=152.23$
Tom	59	170K	1	No	$\text{sqrt} [(59-37)^2+(170-50)^2+(1-2)^2]=122$
Anne	25	40K	4	Yes	$\text{sqrt} [(25-37)^2+(40-50)^2+(4-2)^2]=15.74$
John	37	50K	2	YES	

# How to choose $k$ ?


- If  $k$  is too small it is sensitive to noise points.
- Larger  $k$  works well, but too large  $k$  may include majority points from other classes.



- Rule of thumb is  $k < \sqrt{n}$ ,  $n$  is number of examples.


# Example

**1 NN**




every example in the blue shaded area will be misclassified as the **blue** class

**3 NN**



every example in the blue shaded area will be classified correctly as the **red** class

# Example



- $k$  acts as a smother

# *k*-NN Feature Weighting

- Scale each feature by its importance for classification

$$D(X, Y) = \sqrt{\sum_{i=1}^p w_i (x_i - y_i)^2}$$

1. Can use our **prior knowledge** about which features are more important
2. Can learn the weights  $w_k$  using **cross-validation**

# Feature Normalization

- Distance between neighbors could be dominated by some attributes with relatively large numbers.
  - e.g., income of customers in our previous example.

$$X_s = \frac{X - \min X}{\max X - \min X}$$

- Arises when two features are in different scales.
- Important to normalize those features.
  - Mapping values to numbers between 0 – 1.

# Nominal/Categorical Data

- Distance works naturally with numerical attributes.
- Binary value categorical data attributes can be regarded as 1 or 0.

$$D_H(X, Y) = \sum_{i=1}^k |x_i - y_i|$$


$$x = y \Rightarrow D = 0$$

$$x \neq y \Rightarrow D = 1$$

No.	<i>X</i>	<i>Y</i>	Distance
1	Male	Male	0
2	Male	Female	1



# *k*-NN Classification



# k-NN Classification – Distance

Age	Loan	Default	Distance
25	\$40,000	N	102000
35	\$60,000	N	82000
45	\$80,000	N	62000
20	\$20,000	N	122000
35	\$120,000	N	22000
52	\$18,000	N	124000
23	\$95,000	Y	47000
40	\$62,000	Y	80000
60	\$100,000	Y	42000
48	\$220,000	Y	78000
33	\$150,000	Y	<b>8000</b>
<b>48</b>	<b>\$142,000</b>	<b>?</b>	

Euclidean  
Distance

$$D = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$$

# k-NN Classification – Standardized Distance

Age	Loan	Default	Distance
0.125	0.11	N	0.7652
0.375	0.21	N	0.5200
0.625	0.31	N	<b>0.3160</b>
0	0.01	N	0.9245
0.375	0.50	N	0.3428
0.8	0.00	N	0.6220
0.075	0.38	Y	0.6669
0.5	0.22	Y	0.4437
1	0.41	Y	0.3650
0.7	1.00	Y	0.3861
0.325	0.65	Y	0.3771
<b>0.7</b>	<b>0.61</b>	<b>?</b>	

Standardized  
Variable


$$X_s = \frac{X - \min X}{\max X - \min X}$$

# Strengths and Weaknesses of $k$ -NN

- Strengths
  - Very simple and intuitive.
  - Can be applied to the data from any distribution.
  - Good classification if the number of samples is large enough.
  - Training is very fast
  - Learn complex target functions
  - Do not lose information
- Weaknesses
  - Takes more time to classify a new example.
    - need to calculate and compare distance from new example to all other examples.
  - Choosing  $k$  may be tricky.
  - Need large number of samples for accuracy.
  - Easily fooled by irrelevant attributes

# *k*-NN

- Estimate conditional probability  $\Pr(y|\mathbf{x})$ 
  - Count of data points in class  $y$  in the neighborhood of  $\mathbf{x}$
- Bias and variance tradeoff
  - A small neighborhood  $\rightarrow$  large variance  $\rightarrow$  unreliable estimation
  - A large neighborhood  $\rightarrow$  large bias  $\rightarrow$  inaccurate estimation



# Distance-Weighted $k$ -NN

- Weight the contribution of each close neighbor based on their distances
- Weight function

$$w(\mathbf{x}, \mathbf{x}_i) = \frac{\exp(-|\mathbf{x} - \mathbf{x}_i|_2^2)}{\sum_{j \in \phi_x}^k \exp(-|\mathbf{x} - \mathbf{x}_j|_2^2)}$$

where all  $i \in \phi_x$

- Then,

$$\sum_{j \in \phi_x}^k w(\mathbf{x}, \mathbf{x}_j) = 1$$

- Prediction of  $\mathbf{x}_{new}$

$$y_{new} = \sum_{j \in \phi_{x_{new}}}^k w(\mathbf{x}_{new}, \mathbf{x}_j) y_j$$

# Distance-Weighted $k$ -NN

- Another notation

$$\hat{f}(x_{new}) = \frac{\sum_{j \in \phi_{x_{new}}}^k w_j f(x_j)}{\sum_{j \in \phi_{x_{new}}}^k w_j}$$

where  $w_j = \frac{1}{|x_{new} - x_j|_2^2}$

- Other variants
  - Different weight functions
  - Variable-weighted  $k$ -NN

# Nonparametric Methods

- Parametric distribution models are restricted to specific forms, which may not always be suitable; for example, consider modelling a multimodal distribution with a single, unimodal model.
- Nonparametric approaches make few assumptions about the overall shape of the distribution being modelled.



# Curse of Dimensionality

- Imagine instances described by 20 attributes, but only 2 are relevant to target function
- Curse of dimensionality: nearest neighbor is easily misled when high dimensional  $X$ 
  - “Neighborhood” grows apart.
- Consider  $N$  data points uniformly distributed in a  $p$ -dimensional unit ball centered at origin.
- Consider the  $NN$  estimate at the origin. The mean distance from the origin to the closest data point is:



$$d(p, N) = \left(1 - 2^{-\frac{1}{N}}\right)^{\frac{1}{p}} \approx 1 - \frac{\log N}{p}$$

Eg. 1)  $N=10$ ,  $p=2$ ,  $d = 0.5$

Eg. 2)  $N=10$ ,  $p=50$ ,  $d = 0.98$

# Curse of Dimensionality

- Assume 5000 points uniformly distributed in the unit hypercube and we want to apply 5-NN.
- Suppose our query point is to capture 5 nearest neighbors at the origin ( $p \approx 0.001$ )
  - In 1-dim.: we must go a distance of 0.001 on the average
  - In 2-dim.: we must go  $\sqrt{0.001}$  to get a square that contains 0.001 of the volume
  - In  $d$ -dim.: we must go  $(0.001)^{1/d}$



# ***k*-NN**

- Given a data set with  $N_l$  data points from class  $C_l$  and  $\sum_l N_l = N$ , we have

$$p(\mathbf{x}) = \frac{L}{NV}$$

where  $L$  is the number of data points of  $\mathbf{x}$ , and  $V$  is the total volume of the region ( $\int_R p(\mathbf{x})d\mathbf{x} \approx p(\mathbf{x})V$ ).

- Correspondingly

$$p(\mathbf{x}|C_l) = \frac{L_l}{N_l V}$$

where  $L_l$  is the number of data points of  $\mathbf{x}$  from class  $C_l$ .

- Since  $p(C_l) = N_l/N$ , Bayes' theorem gives

$$p(C_l|\mathbf{x}) = \frac{p(\mathbf{x}|C_l)p(C_l)}{p(\mathbf{x})} = \frac{L_l}{L}$$

# Appendix

- A class label corresponds to a set of points which belong to some region in the feature space  $R$ . If you draw sample points from the actual probability distribution,  $p(\mathbf{x})$ , independently, then the probability of drawing a sample from that class is,

$$P = \int_R p(\mathbf{x}) d\mathbf{x}$$

- What if you have  $N$  points? The probability that  $L$  points of those  $N$  points fall in the region  $R$  follows the binomial distribution,

$$Prob(L) = \binom{N}{L} P^L (1 - P)^{N-L}$$

# Appendix

- As  $N \rightarrow \infty$  this distribution is sharply peaked, so that the probability can be approximated by its mean value  $L/N$ . An additional approximation is that the probability distribution over  $R$  remains approximately constant, so that one can approximate the integral by,

$$P = \int_R p(\mathbf{x}) d\mathbf{x} \approx p(\mathbf{x})V$$

where  $V$  is the total volume of the region. Under this approximations  $p(\mathbf{x}) \approx \frac{L}{NV}$ .

# Appendix

- Now, if we had several classes, we could repeat the same analysis for each one, which would give us,

$$p(x|C_l) = \frac{L_l}{N_l V}$$

- where  $L_l$  is the amount of points from class  $C_l$  which falls within that region and  $N_l$  is the total number of points which fall in that region. Notice  $\sum_l N_l = N$ .
- Repeating the analysis with the binomial distribution, it is easy to see that we can estimate the prior  $p(C_l) = N_l/N$ .
- Using Bayes rule,

$$p(C_l|\mathbf{x}) = \frac{p(\mathbf{x}|C_l)p(C_l)}{p(\mathbf{x})} = \frac{L_l}{L}$$