


Recommender Systems

Instructor: Junghye Lee

Department of Industrial Engineering

junghyelee@unist.ac.kr



Recommender Systems: An Introduction

by [Dietmar Jannach](#), [Markus Zanker](#), [Alexander Felfernig](#), [Gerhard Friedrich](#)

AVERAGE CUSTOMER RATING:

☆☆☆☆☆ ([Be the first to review](#))



[f](#) Registrieren, um sehen zu können, was deinen Freunden gefällt.

FORMAT:

Hardcover

NOOKbook (eBook) - not available

[Tell the publisher you want this in NOOKbook format](#)

NEW FROM BN.COM

~~\$65.00~~ List Price

\$52.00 Online Price
(You Save 20%)

Add to Cart

NEW & USED FROM OUR

New starting at **\$56.46**(You Save 12%)

Used starting at **\$51.98**(You Save 18%)

See All Prices

[Table of Contents](#)

Customers who bought this also bought



Contents

- **What are recommender systems for?**
 - Introduction
- **How do they work (Part I) ?**
 - Collaborative Filtering
- **How do they work (Part II) ?**
 - Content-based Filtering
 - Knowledge-Based Recommendations
- **How to measure their success?**
 - Evaluation techniques

Introduction


Problem domain

- **Recommendation systems (RS) help to match users with items**
 - Ease information overload
 - Sales assistance (guidance, advisory, persuasion,...)

RS are software agents that elicit the interests and preferences of individual consumers [...] and make recommendations accordingly.

They have the potential to support and improve the quality of the decisions consumers make while searching for and selecting products online.

» [Xiao & Benbasat, MISQ, 2007]



- **Different system designs / paradigms**
 - Based on availability of exploitable data
 - Implicit and explicit user feedback
 - Domain characteristics



Purpose and success criteria (1)

- **Different perspectives/aspects**
 - Depends on domain and purpose
 - No holistic evaluation scenario exists

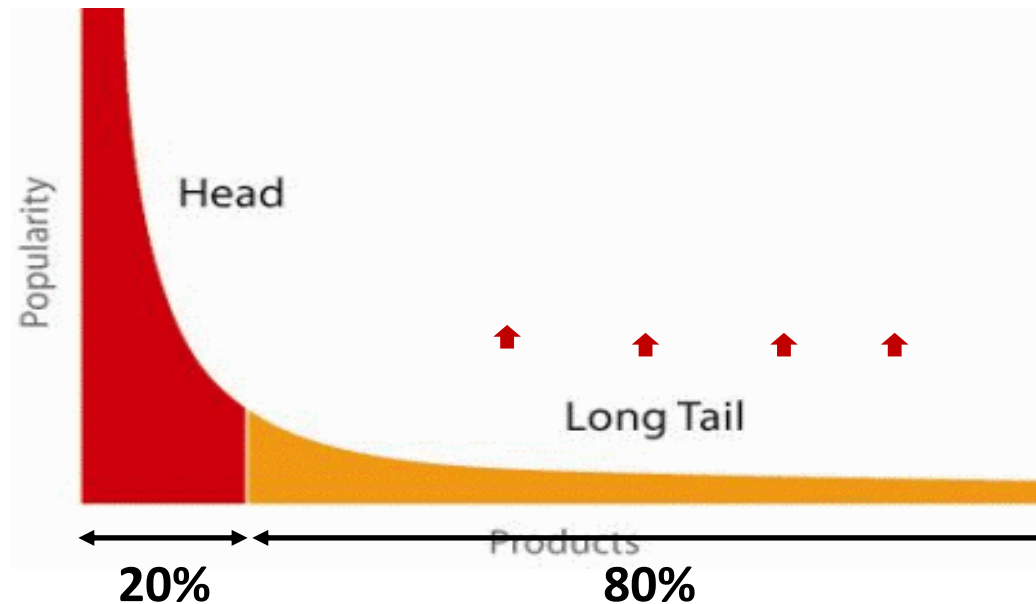
- **Retrieval perspective**
 - Reduce search costs
 - Provide "correct" proposals
 - Users know in advance what they want

- **Recommendation perspective**
 - Serendipity – identify items from the Long Tail
 - Users did not know about existence

When does a RS do its job well?

**"Recommend widely unknown items
that users might actually like!"**

Recommend items from the long tail



**20% of items
accumulate 74% of all
positive ratings**

Items rated > 3 in MovieLens 100K dataset

Purpose and success criteria (2)

- **Prediction perspective**

- Predict to what degree users like an item
- Most popular evaluation scenario in research

- **Interaction perspective**

- Give users a "good feeling"
- Educate users about the product domain
- Convince/persuade users - explain

- **Finally, conversion perspective**

- Commercial situations
- Increase "hit", "clickthrough", "lookers to bookers" rates
- Optimize sales margins and profit

Recommender systems


- **RS seen as a function**
 - **Given:**
 - User model (e.g. ratings, preferences, demographics, situational context)
 - Items (with or without description of item characteristics)
 - **Find:**
 - Relevance score. Used for ranking.
-
- **Relation to Information Retrieval:**
 - IR is finding material [..] of an unstructured nature [..] that satisfies an information need from within large collections [..].
 - » [Manning et al., CUP, 2008]

Recommender systems

- **RS seen as a function**
 - **Given:**
 - User model (e.g. ratings, preferences, demographics, situational context)
 - Items (with or without description of item characteristics)
 - **Find:**
 - Relevance score. Used for ranking.
-
- **Relation to Information Retrieval:**
 - IR is finding material [..] of an unstructured nature [..] that satisfies an information need from within large collections [..].
 - » [Manning et al., CUP, 2008]

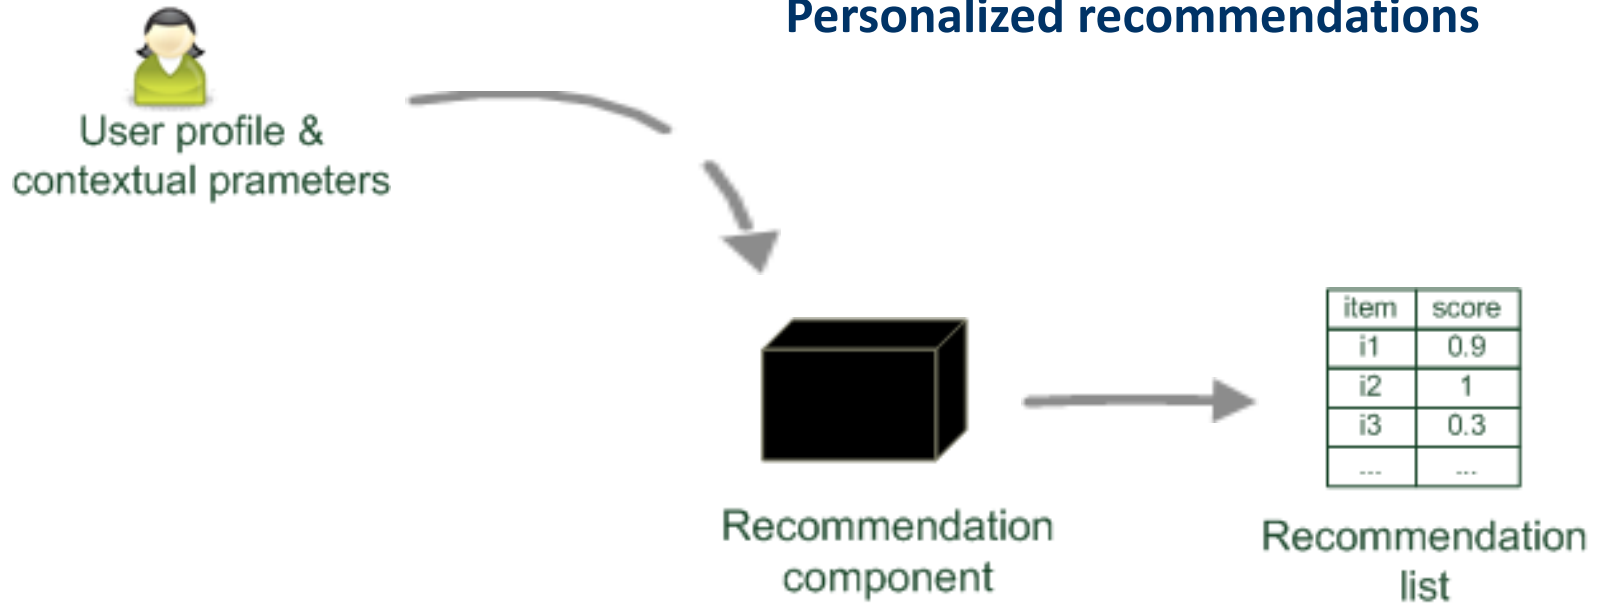
Paradigms of recommender systems

Recommender systems reduce information overload by estimating relevance

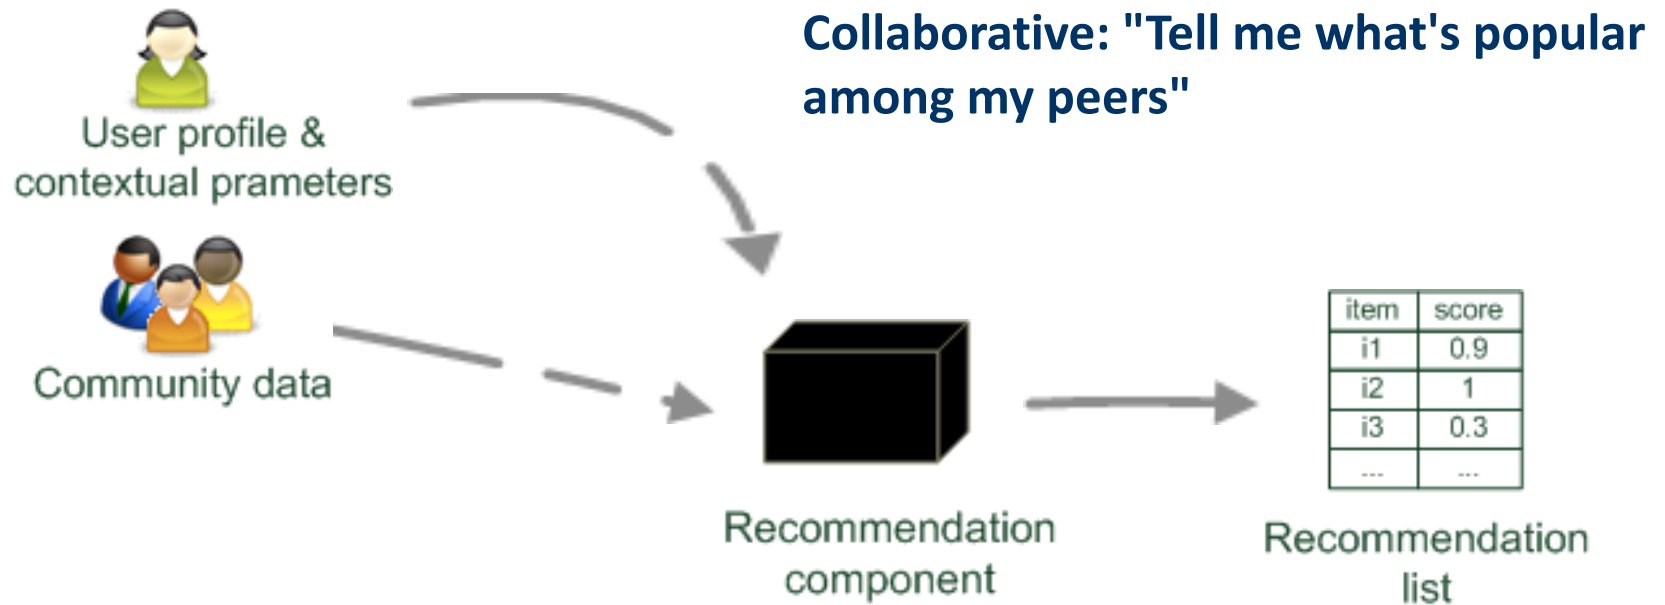


Paradigms of recommender systems

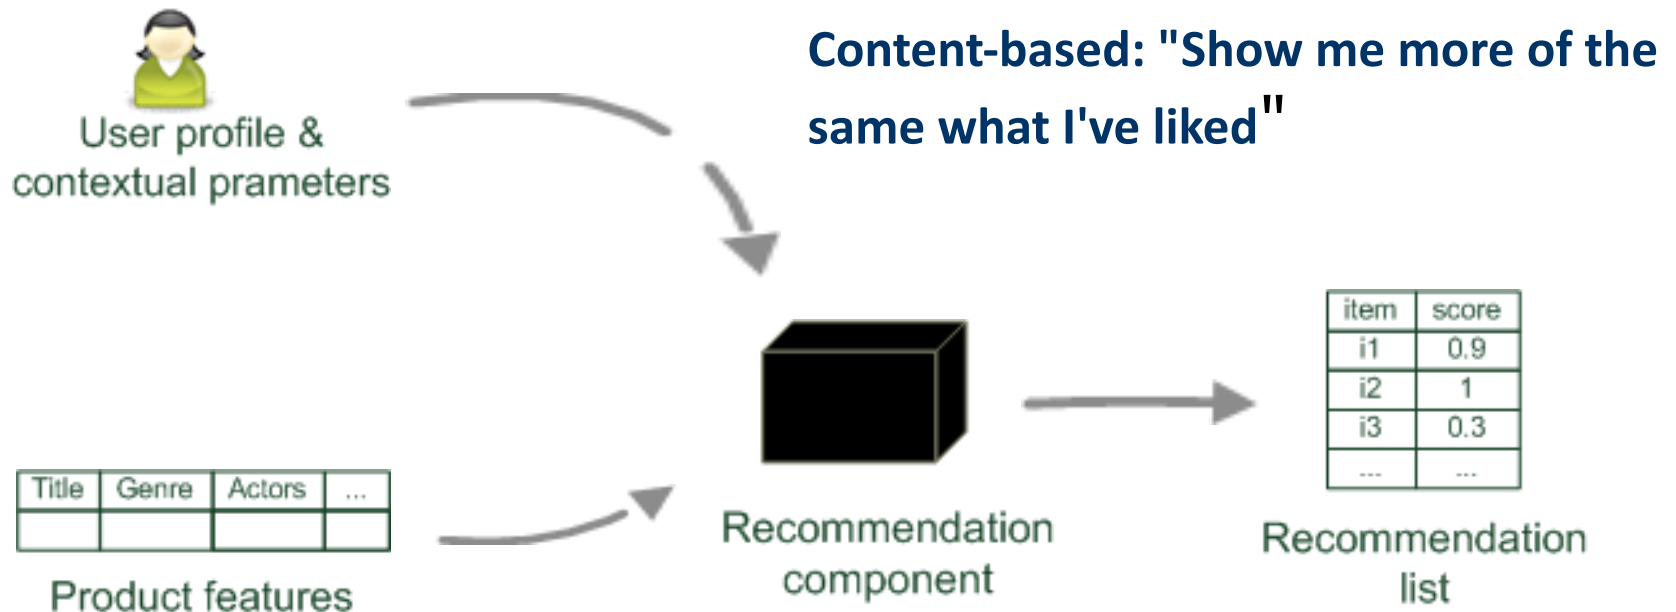
Personalized recommendations



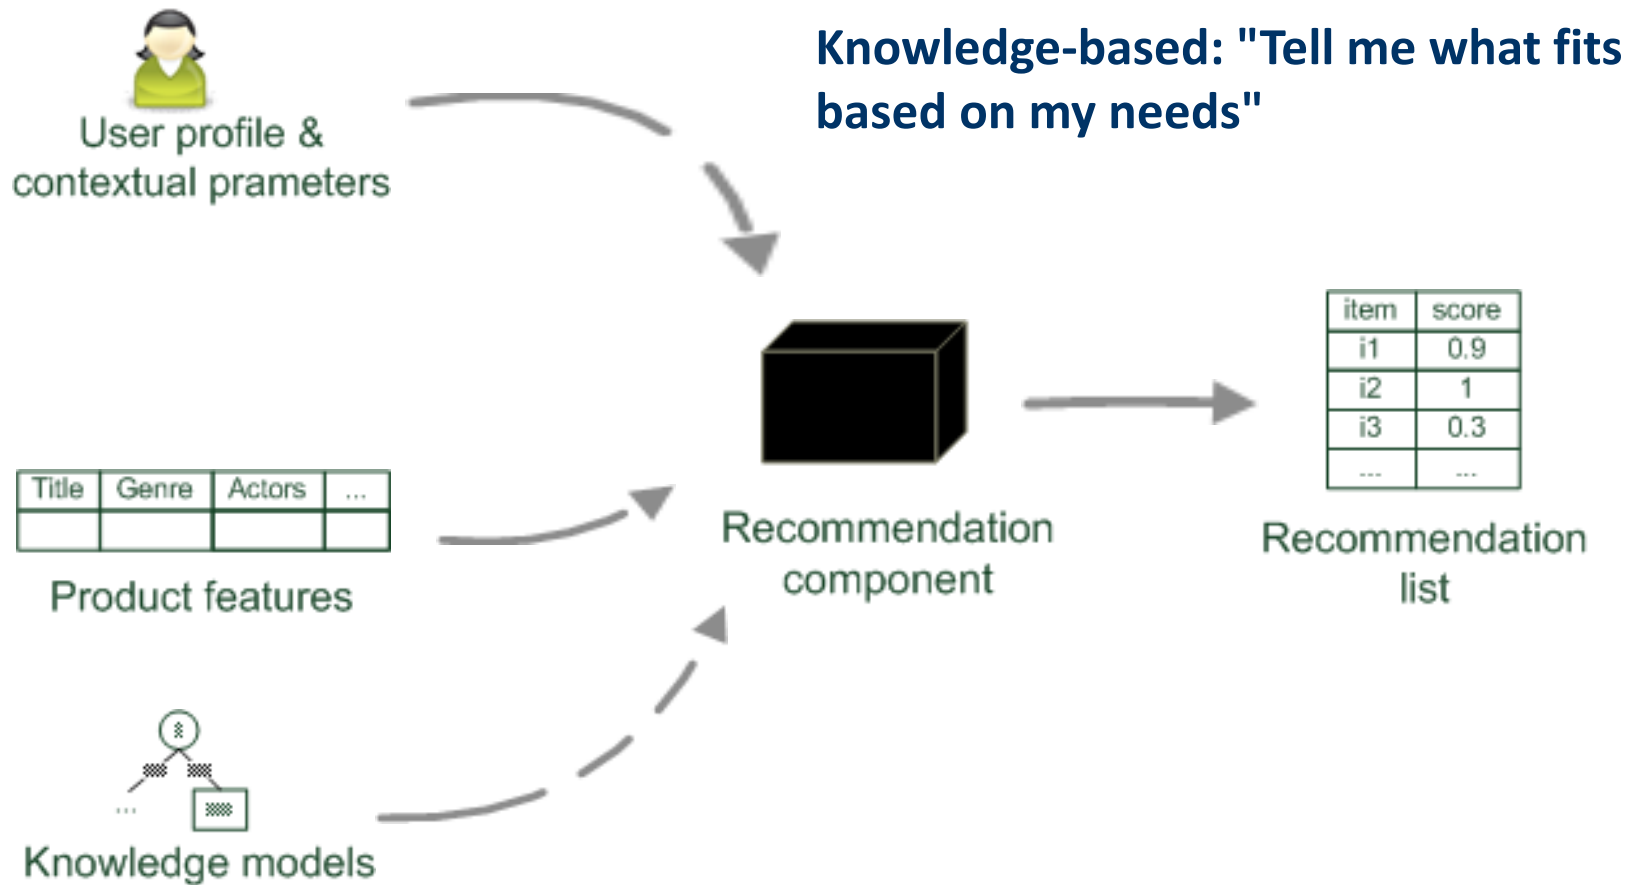
Paradigms of recommender systems



Paradigms of recommender systems



Paradigms of recommender systems




Collaborative Filtering

Collaborative Filtering (CF)

- **The most prominent approach to generate recommendations**
 - used by large, commercial e-commerce sites
 - well-understood, various algorithms and variations exist
 - applicable in many domains (book, movies, DVDs, ..)
- **Approach**
 - use the "wisdom of the crowd" to recommend items
- **Basic assumption and idea**
 - Users give ratings to catalog items (implicitly or explicitly)
 - Customers who had similar tastes in the past, will have similar tastes in the future

User-based vs. Item-based CF

- **The way to recommend either item 3 and item 4 to user**
 - user is similar to user 1, so recommend item 3 which is preferred by user 1
 - In terms of USER
 - user prefers item 1 and item 1 is similar to item 3, so recommend item 3
 - In terms of ITEM



User-based nearest-neighbor collaborative filtering (1)

- **The basic technique:**

- Given an "active user" (Alice) and an item p not yet seen by Alice
- The *goal is to estimate Alice's rating for this item*, e.g., by
 - find a set of users (peers) who liked the same items as Alice in the past and who have rated item p
 - use, e.g. the average of their ratings to predict, if Alice will like item p
 - do this for all items Alice has not seen and recommend the best-rated

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

User-based nearest-neighbor collaborative filtering (2)

■ Some first questions

- How do we measure similarity?
- How many neighbors should we consider?
- How do we generate a prediction from the neighbors' ratings?

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

Measuring user similarity

- A popular similarity measure in user-based CF

$$sim(a, b) = \frac{\sum_{p \in P} (r_{a,p} - \bar{r}_a) (r_{b,p} - \bar{r}_b)}{\sqrt{\sum_{p \in P} (r_{a,p} - \bar{r}_a)^2} \sqrt{\sum_{p \in P} (r_{b,p} - \bar{r}_b)^2}}$$


a, b : users

$r_{a,p}$: rating of user a for item p

P : set of items, rated both by a and b

\bar{r}_a, \bar{r}_b : user's average ratings


	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1



sim = 0.85
sim = 0.70
sim = -0.79

Pearson correlation

- Takes differences in rating behavior into account



- Works well in usual domains, compared with alternative measures
 - such as cosine similarity

Making predictions

- A common prediction function:

$$pred(a, p) = \bar{r}_a + \frac{\sum_{b \in N} sim(a, b) * (r_{b,p} - \bar{r}_b)}{\sum_{b \in N} |sim(a, b)|}$$

where user a and item p , N : the number of users

- Calculate, whether the neighbors' ratings for the unseen item p are higher or lower than their average
- Combine the rating differences – use the similarity with as a weight
- Add/subtract the neighbors' bias from the active user's average and use this as a prediction

Improving the metrics / prediction function

- **Not all neighbor ratings might be equally "valuable"**
 - Agreement on commonly liked items is not so informative as agreement on controversial items
 - **Possible solution:** Give more weight to items that have a higher variance
- **Value of number of co-rated items**
 - Use "significance weighting", by e.g., linearly reducing the weight when the number of co-rated items is low
- **Case amplification**
 - Intuition: Give more weight to "very similar" neighbors, i.e., where the similarity value is close to 1.
- **Neighborhood selection**
 - Use similarity threshold or fixed number of neighbors

Memory-based and model-based approaches

- **User-based CF is said to be "memory-based"**
 - the rating matrix is directly used to find neighbors / make predictions
 - does not scale for most real-world scenarios
 - large e-commerce sites have tens of millions of customers and millions of items
- **Model-based approaches**
 - based on an offline pre-processing or "model-learning" phase
 - at run-time, only the learned model is used to make predictions
 - models are updated / re-trained periodically
 - large variety of techniques used
 - model-building and updating can be computationally expensive

2001: *Item-based collaborative filtering recommendation algorithms*, B. Sarwar et al., WWW 2001

- **Scalability issues arise with U2U if many more users than items ($m \gg n$, $m = |\text{users}|$, $n = |\text{items}|$)**
 - e.g. amazon.com
 - Space complexity $O(m^2)$ when pre-computed
 - Time complexity for computing Pearson $O(m^2n)$

- **High sparsity leads to few common ratings between two users**

- **Basic idea: "Item-based CF exploits relationships between items first, instead of relationships between users"**

Item-based collaborative filtering

- **Basic idea:**
 - Use the similarity between items (and not users) to make predictions
- **Example:**
 - Look for items that are similar to Item5
 - Take Alice's ratings for these items to predict the rating for Item5

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

Similarity measure

- **Produces better results in item-to-item filtering**
 - for some datasets, no consistent picture in literature
- **Ratings are seen as vector in n-dimensional space**
- **Cosine similarity**

$$\text{sim}(\vec{p}, \vec{q}) = \frac{\vec{p} \cdot \vec{q}}{|\vec{p}| * |\vec{q}|}$$

- **Cosine similarity**
- **Adjusted cosine similarity**
 - take average user ratings into account, transform the original ratings
 - U: set of users who have rated both items p and q

$$\text{sim}(p, q) = \frac{\sum_{u \in U} (r_{u,p} - \bar{r}_p) (r_{u,q} - \bar{r}_q)}{\sqrt{\sum_{u \in U} (r_{u,p} - \bar{r}_p)^2} \sqrt{\sum_{u \in U} (r_{u,q} - \bar{r}_q)^2}}$$

Similarity measure

$$pred(a, p) = \frac{\sum_{q \in S(p, a, k)} sim(p, q) * r_{a, q}}{\sum_{q \in S(p, a, k)} |sim(p, q)|}$$

- **Produces better results in item-to-item filtering**
 - for some datasets, no consistent picture in literature
- **Ratings are seen as vector in n-dimensional space**
- **Cosine similarity**

$$sim(\vec{p}, \vec{q}) = \frac{\vec{p} \cdot \vec{q}}{|\vec{p}| * |\vec{q}|}$$

- **Cosine similarity**
- **Adjusted cosine similarity**
 - take average user ratings into account, transform the original ratings
 - U: set of users who have rated both items p and q

$$sim(p, q) = \frac{\sum_{u \in U} (r_{u, p} - \bar{r}_p) (r_{u, q} - \bar{r}_q)}{\sqrt{\sum_{u \in U} (r_{u, p} - \bar{r}_p)^2} \sqrt{\sum_{u \in U} (r_{u, q} - \bar{r}_q)^2}}$$

Pre-processing for item-based filtering

- **Item-based filtering does not solve the scalability problem itself**
- **Pre-processing approach by Amazon.com (in 2003)**
 - Calculate all pair-wise item similarities in advance
 - The neighborhood to be used at run-time is typically rather small, because only items are taken into account which the user has rated
 - Item similarities are supposed to be more stable than user similarities
- **Memory requirements**
 - Up to n^2 pair-wise similarities to be memorized (n = number of items) in theory
 - In practice, this is significantly lower (items with no co-ratings)
 - Further reductions possible
 - Minimum threshold for co-ratings (items, which are rated at least by l users)
 - Limit the size of the neighborhood (k might affect recommendation accuracy)

More on ratings

- **Pure CF-based systems only rely on the rating matrix**
- **Explicit ratings**
 - Most commonly used (1 to 5, 1 to 7 Likert response scales)
 - Research topics
 - "Optimal" granularity of scale; indication that 10-point scale is better accepted in movie domain
 - Multidimensional ratings (multiple ratings per movie)
 - Challenge
 - Users not always willing to rate many items; sparse rating matrices
 - How to stimulate users to rate more items?
- **Implicit ratings**
 - clicks, page views, time spent on some page, demo downloads, and etc. can be used in addition to explicit ones
 - question of correctness of interpretation

Data sparsity problems

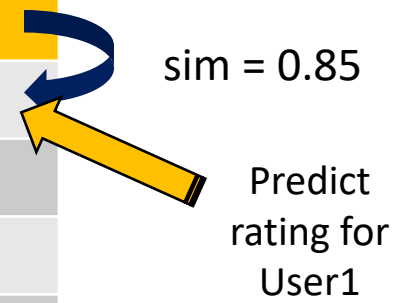
- **Cold start problem**
 - How to recommend new items? What to recommend to new users?
- **Straightforward approaches**
 - Ask/force users to rate a set of items
 - Use another method (e.g., content-based, demographic or simply non-personalized) in the initial phase
- **Alternatives**
 - Use better algorithms (beyond nearest-neighbor approaches)
 - Example:
 - In nearest-neighbor approaches, the set of sufficiently similar neighbors might be too small to make good predictions

Example algorithms for sparse datasets

■ Recursive CF

- Assume there is a very close neighbor b of a who however has not rated the target item p yet.
- Idea:
 - Apply CF-method recursively and predict a rating for item p for the neighbor
 - Use this predicted rating instead of the rating of a more distant direct neighbor

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	?
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1




sim = 0.85

Predict rating for User1

Graph-based methods

- "Spreading activation" (sketch)
 - Idea: Use paths of lengths > 3 to recommend items
 - Length 3: Recommend Item3 to User1
 - Length 5: Item1 also recommendable



**Bipartite graph,
odd number!**

More model-based approaches

- **Plethora of different techniques proposed in the last years, e.g.,**
 - Matrix factorization techniques, statistics
 - singular value decomposition, principal component analysis
 - Association rule mining
 - compare: shopping basket analysis
 - Probabilistic models
 - clustering models, Bayesian networks, probabilistic Latent Semantic Analysis
 - Various other machine learning approaches
- **Costs of pre-processing**
 - Usually not discussed
 - Incremental updates possible?






2000: *Application of Dimensionality Reduction in Recommender System*, B. Sarwar et al., WebKDD Workshop

- **Basic idea: Trade more complex offline model building for faster online prediction generation**
- **Singular Value Decomposition for dimensionality reduction of rating matrices**
 - Captures important factors/aspects and their weights in the data
 - factors can be genre, actors but also non-understandable ones
 - Assumption that k dimensions capture the signals and filter out noise ($K = 20$ to 100)
- **Constant time to make recommendations**
- **Approach also popular in IR (Latent Semantic Indexing), data compression,...**

Matrix factorization

- SVD: $M_k = U_k \times \Sigma_k \times V_k^T$


U_k	Dim1	Dim2
Alice	0.47	-0.30
Bob	-0.44	0.23
Mary	0.70	-0.06
Sue	0.31	0.93

V_k^T					
Dim1	-0.44	-0.57	0.06	0.38	0.57
Dim2	0.58	-0.66	0.26	0.18	-0.36

Σ_k	Dim1	Dim2
Dim1	5.63	0
Dim2	0	3.23

- Prediction: $\hat{r}_{ui} = \bar{r}_u + U_k(\text{Alice}) \times \Sigma_k \times V_k^T(\text{EPL})$
 $= 3 + 0.84 = \mathbf{3.84}$

A picture says ...



Association rule mining

- **Commonly used for shopping behavior analysis**

- aims at detection of rules such as

"If a customer purchases baby-food then he also buys diapers in 70% of the cases"

- **Association rule mining algorithms**

- can detect rules of the form $X \Rightarrow Y$ (e.g., baby-food \Rightarrow diapers) from a set of sales transactions $D = \{t_1, t_2, \dots, t_n\}$
 - measure of quality: support, confidence
 - used e.g. as a threshold to cut off unimportant rules

- let $\sigma(X) = \frac{|\{x | x \subseteq t_1, t_1 \in D\}|}{|D|}$

- support = $\frac{\sigma(X \cup Y)}{|D|}$, confidence = $\frac{\sigma(X \cup Y)}{\sigma(X)}$

$$support = \frac{|X \cup Y|}{|Transactions|}$$

$$confidence = \frac{|X \cup Y|}{|X|}$$

Recommendation based on Association Rule Mining

- **Simplest approach**

- transform 5-point ratings into binary ratings (1 = above user average)

- **Mine rules such as**

- Item1 => Item5
 - support (2/4), confidence (2/2) (without Alice)

- **Make recommendations for Alice (basic method)**

- Determine "relevant" rules based on Alice's transactions (the above rule will be relevant as Alice bought Item1)
- Determine items not already bought by Alice
- Sort the items based on the rules' confidence values

- **Different variations possible**

- dislike statements, user associations ..

	Item1	Item2	Item3	Item4	Item5
Alice	1	0	0	0	?
User1	1	0	1	0	1
User2	1	0	1	0	1
User3	0	0	0	1	1
User4	0	1	1	0	0

Probabilistic methods

- **Basic idea (simplistic version for illustration):**

- given the user/item rating matrix
- determine the probability that user Alice will like an item p
- base the recommendation on such these probabilities

- **Calculation of rating probabilities based on Bayes Theorem**

- How probable is rating value "1" for Item5 given Alice's previous ratings?
- Corresponds to conditional probability $P(\text{Item5}=1 | X)$, where
 - $X = \text{Alice's previous ratings} = (\text{Item1}=1, \text{Item2}=3, \text{Item3}= \dots)$
- Can be estimated based on Bayes Theorem

$$P(Y|X) = \frac{P(X|Y) \times P(Y)}{P(X)} \quad P(Y|X) = \frac{\prod_{i=1}^d P(X_i|Y) \times P(Y)}{P(X)}$$

- Assumption: Ratings are independent, d = number of attributes in X , # of items

Calculation of probabilities in simplistic approach

	Item1	Item2	Item3	Item4	Item5
Alice	1	3	3	2	?
User1	2	4	2	2	4
User2	1	3	3	5	1
User3	4	5	2	3	3
User4	1	1	5	2	1

$X = (\text{Item1} = 1, \text{Item2} = 3, \text{Item3} = \dots)$

$$\begin{aligned} P(X|\text{Item5}=1) &= P(\text{Item1}=1|\text{Item5}=1) \times P(\text{Item2}=3|\text{Item5}=1) \times \\ &\quad P(\text{Item3}=3|\text{Item5}=1) \times P(\text{Item4}=2|\text{Item5}=1) \\ &= 2/4 \times 1/4 \times 1/4 \times 1/4 \\ &\approx 0.0078125 \end{aligned}$$

$$\begin{aligned} P(X|\text{Item5}=2) &= P(\text{Item1}=1|\text{Item5}=2) \times P(\text{Item2}=3|\text{Item5}=2) \times \\ &\quad P(\text{Item3}=3|\text{Item5}=2) \times P(\text{Item4}=2|\text{Item5}=2) \\ &= 0/4 \times \dots \times \dots \times \dots \\ &= 0 \end{aligned}$$



■ More to consider

- Zeros (smoothing required)
- like/dislike simplification possible

Practical probabilistic approaches

- **Use a cluster-based approach**
 - assume users fall in a small number of subgroups (clusters)
 - Make predictions based on estimates
 - probability of Alice falling into cluster c
 - probability of Alice liking item i given a certain cluster and her previous ratings
 - Based on model-based clustering (mixture model)
 - Number of classes and model parameters have to be learned from data in advance (EM algorithm)
- **Others:**
 - Bayesian Networks, Probabilistic Latent Semantic Analysis,
- **Empirical analysis shows:**
 - Probabilistic methods lead to relatively good results (movie domain)
 - No consistent winner; small memory-footprint of network model

Collaborative Filtering Issues

- **Pros:** 
 - well-understood, works well in some domains, no knowledge engineering required
- **Cons:** 
 - requires user community, sparsity problems, no integration of other knowledge sources, no explanation of results
- **What is the best CF method?**
 - In which situation and which domain? Inconsistent findings; always the same domains and data sets; differences between methods are often very small (1/100)
- **How to evaluate the prediction quality?**
 - MAE / RMSE: What does an MAE of 0.7 actually mean?
 - Serendipity: Not yet fully understood
- **What about multi-dimensional ratings?**